

Paras Malik

Github|paras.sde@gmail.com|+917204537001|Bangalore,IN

Machine Learning, Distributed Systems, Large Scale Systems, Big Data, Tech Leader

EXPERIENCE

Google Staff Software Engineer

March 2022 - Present|Bangalore

- Spearheaded the comprehension of B2B queries by creating query classification and annotation models and seamlessly incorporating them into the Ads serving process
- Initiated a project to identify deficiencies in Ads display images and put forth innovative solutions to enhance their quality, resulting in an estimated 3% increase in CTR for targeted product listing ads.
- Guided junior engineers through problem-solving by clearly defining project scope and providing support to overcome obstacles.

Media.Net Sr. EM|Sr. Tech Lead|Sr. SDE

Jan 2014 - March 2022|Bangalore

- Managed a multi discipline organization(Engineering, Data Science, Product) with over 20 people and 4 Leads across 2 different products and 2 locations
- Responsible for building Keyword intelligence and Bidder for DSP platforms
- Created realtime bidder that bid on 15 Billion requests, serving 100s of millions of impressions each day and generating 85 Million USD of annual revenue and 30 Million USD of annual profit for the company
- My team consists of 20 brilliant SE/SRE/DS/PM/DB engineers who deliver quality projects with focus on performance, reliability and business objectives
- Responsible for Appraisals, Promotion recommendations, Ideation, Engineering, Product & Project Management
- Actively involved in hiring and mentoring

Amazon.com Software Development Engineer

Oct 2010 - Dec 2013|Bangalore

- Developed Java services for Kindle registration and book lending
- Developed Amazon Instant Video Living room client applications using Java

Directi Software Development Engineer

June 2010 - Oct 2010|Mumbai

- Developed Implemented caching for java application using JNI

SKILLS

Programming/Scripting Languages:Java, Golang, Python, C/C++, Bash Scripting, R, Javascript

Databases/Caching Technologies:Redis, HBase, Aerospike, Mysql, Mssql, Memcache, Cassandra

Big Data Solutions:Hive, Hadoop, Spark/Py-spark/Spark Structured Streaming

Machine Learning Techniques:Statistical Modeling, UCB, XGBoost, LightGBM, Random Forest, FM, FFM, LDA, NaiveBayesClassifier

Messaging Systems:Kafka

Search Technologies:ES, Lucene

CICD:Docker, Messos, Kubernetes, Jenkins

Management Skills:Product & Project Management, Hiring, Team building & scaling, Operational excellence

Other Technologies:Tomcat, SOA (Service Discovery, Client side load balancing), AWS(Auto-scaling/S3/ELB/EBS/Machine-Selection/Placement Groups)

EDUCATION

Masters in Computer Application Indian Institute of Technology|2010

PROJECTS

RTB Bidder Service Java|Redis|Aerospike

- Developed & Maintained a complex API to compute and deliver 1/2 Million bids per second for RTB auctions.

Realtime Configuration Manager Mysql|Kafka|Java

- Developed system which update 100s of 1000s of frequently changing configurations to production in realtime

ML Model Computation Service

Golang|S3|Redis|LightGBM|Catboost|Consul|Java

- Developed low latency Golang service to compute various ML model.
- Used Consul for service discovery in autoscaled environment using heterogeneous hosts.
- Developed intelligent Java client with capability of WRR

Event Stream Synchronization & Processing Spark Structured Streaming|Kafka

- Developed structured streaming application to time synchronize multiple streams of different events in different Kafka topics.

Data Pipelines Spark|Hadoop|Oozie|Mysql

- Developed complex data pipelines to process large amount of data

SRE Projects ELK|Autoscaling|AWS|GCP

- Developed Autoscaling algorithm to add/remove hosts
- Migration from AWS to GCP and Messos to GKE

ML Projects

- Developed System to classify web pages as commercial non-commercial using LDA/NaiveBayes/XGBoost
- Developed RPM/CTR Prediction models using LightGBM/Catboost/XGboost/Statistical Modeling
- Developed Slideshow detection system which automatically identify if given page is part of slideshow.
- Worked on UCB technique to explore new web pages

Redis to RedisCluster Migration Redis|RedisCluster|Jedis

- Implemented functionality to handle pipeline requests in Jedis
- Used different functionality to optimize RedisCluster
- Benchmarked and controlled migration

HBase & RedisCluster to Aerospike Migration HBase|Aerospike|RedisCluster

- Defined Data model to enable migration in a cost efficient way.